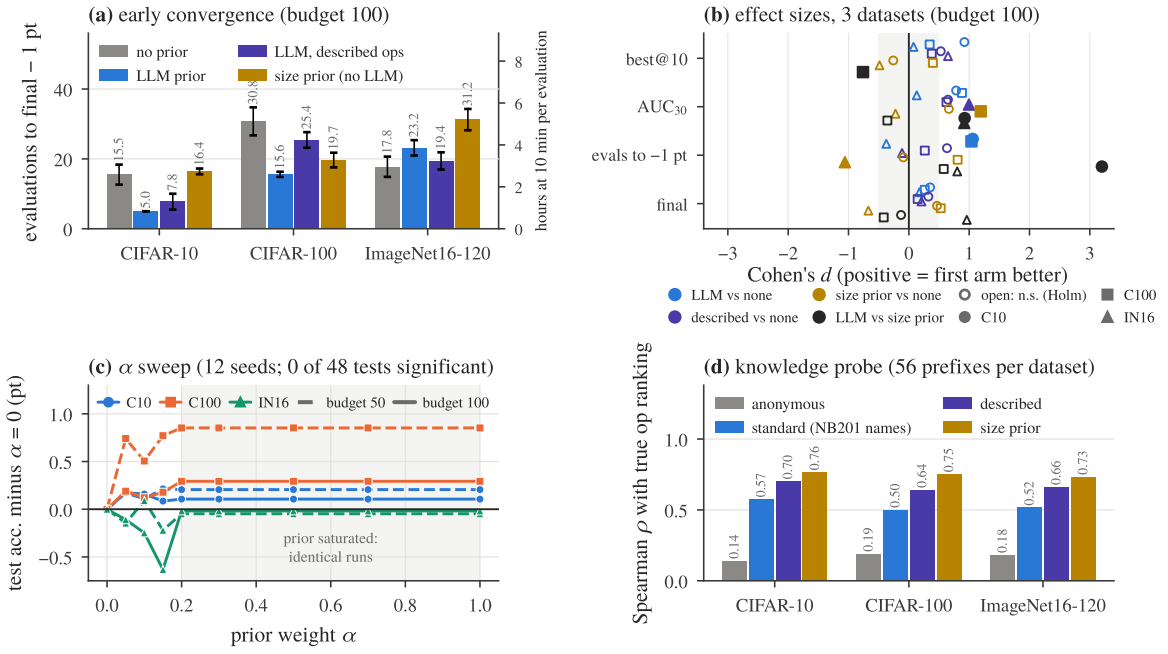
(a) early convergence (budget 100)
no prior
LLM prior
LLM, described ops
size prior (no LLM)
evaluations to final − 1 pt
hours at 10 min per evaluation
5.5
5.0
7.8
16.4
30.8
15.6
25.4
19.7
17.8
23.2
19.4
31.2
CIFAR-10
CIFAR-100
ImageNet16-120
(b) effect sizes, 3 datasets (budget 100)
best@10
AUC30
evals to −1 pt
final
Cohen's d (positive = first arm better)
LLM vs none
described vs none
size prior vs none
LLM vs size prior
open: n.s. (Holm)
C10
C100
IN16
(c) α sweep (12 seeds; 0 of 48 tests significant)
C10
C100
IN16
budget 50
budget 100
test acc. minus α = 0 (pt)
prior saturated:
identical runs
prior weight α
(d) knowledge probe (56 prefixes per dataset)
anonymous
standard (NB201 names)
described
size prior
Spearman ρ with true op ranking
0.14
0.57
0.70
0.76
0.19
0.50
0.64
0.75
0.18
0.52
0.66
0.73
CIFAR-10
CIFAR-100
ImageNet16-120